

PANDAS

DATA ANALYSIS ASSIGNMENT #1

ANALYZING EMPLOYEE ATTRITION



[Link: Pandas Documentation]
(<https://pandas.pydata.org/docs/>)

Files needed for this assignment:

[Employee_Attrition.csv](#)

```
In [2]: # set up notebook to display multiple output in one cell

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

print('The notebook is set up to display multiple output in one cell.')
```

The notebook is set up to display multiple output in one cell.

```
In [3]: import pandas as pd
import numpy as np
```

Note: Throughout this assignment, add cells as needed.

Task #1: Reading the Dataset 1. Read in the `[**Employee_Attrition.csv**]`
(https://drive.google.com/file/d/1E-XLIZdyUYGf-cCvjx6KaNNICX1ATsVo/view?usp=share_link)

dataset and store the results in a DataFrame.

2. Repeat Step #1, but this time set one of the columns as the index.

```
In [4]: employee = pd.read_csv('Employee_Attrition.csv', sep=';', index_col='EmployeeNumber', skip:
employee
```

Out[4]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|--|-----|-----------|----------------|-----------|------------|------------------|-------|
|--|-----|-----------|----------------|-----------|------------|------------------|-------|

| EmployeeNumber | | | | | | | |
|----------------|-----|-----|-------------------|------|------------------------|--|-----|
| 1 | 41 | Yes | Travel_Rarely | 1102 | Sales | | 1 |
| 2 | 49 | No | Travel_Frequently | 279 | Research & Development | | 8 |
| 4 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | | 2 |
| 5 | 33 | No | Travel_Frequently | 1392 | Research & Development | | 3 |
| 7 | 27 | No | Travel_Rarely | 591 | Research & Development | | 2 |
| ... | ... | ... | ... | ... | ... | | ... |
| 2061 | 36 | No | Travel_Frequently | 884 | Research & Development | | 23 |
| 2062 | 39 | No | Travel_Rarely | 613 | Research & Development | | 6 |
| 2064 | 27 | No | Travel_Rarely | 155 | Research & Development | | 4 |
| 2065 | 49 | No | Travel_Frequently | 1023 | Sales | | 2 |
| 2068 | 34 | No | Travel_Rarely | 628 | Research & Development | | 8 |

1470 rows × 34 columns

Task #2: Use DataFrame Attributes to Inspect the Anatomy of the DataFrame 1.

Write code to access the DataFrame's index (i.e., to access the row labels).

2. Write code to access the DataFrame's column names/columns (column index).

3. Write code to determine the data type of each variable in your DataFrame.

4. Write code to access the DataFrame's data (i.e., just the data without the index or column names).

5. Write code to determine the shape (i.e., the dimensions) of the DataFrame.

```
In [5]: employee.index
```

```
Out[5]: Int64Index([ 1, 2, 4, 5, 7, 8, 10, 11, 12, 13,
...,
2054, 2055, 2056, 2057, 2060, 2061, 2062, 2064, 2065, 2068],
dtype='int64', name='EmployeeNumber', length=1470)
```

```
In [6]: employee.columns
```

```
Out[6]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',  
            'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',  
            'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',  
            'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',  
            'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18',  
            'OverTime', 'PercentSalaryHike', 'PerformanceRating',  
            'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',  
            'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',  
            'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',  
            'YearsWithCurrManager'],  
           dtype='object')
```

```
In [7]: employee.dtypes
```

```
Out[7]: Age                int64  
Attrition                object  
BusinessTravel            object  
DailyRate                int64  
Department                object  
DistanceFromHome          int64  
Education                 int64  
EducationField            object  
EmployeeCount             int64  
EnvironmentSatisfaction    int64  
Gender                    object  
HourlyRate                int64  
JobInvolvement            int64  
JobLevel                  int64  
JobRole                    object  
JobSatisfaction            int64  
MaritalStatus             object  
MonthlyIncome             int64  
MonthlyRate               int64  
NumCompaniesWorked        int64  
Over18                    object  
OverTime                  object  
PercentSalaryHike         int64  
PerformanceRating         int64  
RelationshipSatisfaction    int64  
StandardHours             int64  
StockOptionLevel          int64  
TotalWorkingYears         int64  
TrainingTimesLastYear     int64  
WorkLifeBalance           int64  
YearsAtCompany            int64  
YearsInCurrentRole        int64  
YearsSinceLastPromotion   int64  
YearsWithCurrManager      int64  
dtype: object
```

```
In [8]: employee.values
```

```
Out[8]: array([[41, 'Yes', 'Travel_Rarely', ..., 4, 0, 5],
        [49, 'No', 'Travel_Frequently', ..., 7, 1, 7],
        [37, 'Yes', 'Travel_Rarely', ..., 0, 0, 0],
        ...,
        [27, 'No', 'Travel_Rarely', ..., 2, 0, 3],
        [49, 'No', 'Travel_Frequently', ..., 6, 0, 8],
        [34, 'No', 'Travel_Rarely', ..., 3, 1, 2]], dtype=object)
```

```
In [9]: employee.shape
```

```
Out[9]: (1470, 34)
```

Task #3: Use DataFrame Methods to Inspect Your Data 1. Write code to access the first 5 rows of the data.

2. Write code to access the first 8 rows of the data.

3. Write code to access the last 5 rows of the data.

4. Write code to access the last 7 rows of the data.

5. Write code to get detailed information about your DataFrame.

```
In [10]: employee.head()
```

```
Out[10]:
```

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|-----------------------|-----|-----------|-------------------|-----------|------------------------|------------------|-------|
| EmployeeNumber | | | | | | | |
| 1 | 41 | Yes | Travel_Rarely | 1102 | Sales | | 1 |
| 2 | 49 | No | Travel_Frequently | 279 | Research & Development | | 8 |
| 4 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | | 2 |
| 5 | 33 | No | Travel_Frequently | 1392 | Research & Development | | 3 |
| 7 | 27 | No | Travel_Rarely | 591 | Research & Development | | 2 |


5 rows × 34 columns

```
In [11]: employee.head(8)
```

Out[11]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|----------------|-----|-----------|-------------------|-----------|------------------------|------------------|-------|
| EmployeeNumber | | | | | | | |
| 1 | 41 | Yes | Travel_Rarely | 1102 | Sales | 1 | |
| 2 | 49 | No | Travel_Frequently | 279 | Research & Development | 8 | |
| 4 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | 2 | |
| 5 | 33 | No | Travel_Frequently | 1392 | Research & Development | 3 | |
| 7 | 27 | No | Travel_Rarely | 591 | Research & Development | 2 | |
| 8 | 32 | No | Travel_Frequently | 1005 | Research & Development | 2 | |
| 10 | 59 | No | Travel_Rarely | 1324 | Research & Development | 3 | |
| 11 | 30 | No | Travel_Rarely | 1358 | Research & Development | 24 | |

8 rows × 34 columns




In [12]: `employee.tail()`

Out[12]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|----------------|-----|-----------|-------------------|-----------|------------------------|------------------|-------|
| EmployeeNumber | | | | | | | |
| 2061 | 36 | No | Travel_Frequently | 884 | Research & Development | 23 | |
| 2062 | 39 | No | Travel_Rarely | 613 | Research & Development | 6 | |
| 2064 | 27 | No | Travel_Rarely | 155 | Research & Development | 4 | |
| 2065 | 49 | No | Travel_Frequently | 1023 | Sales | 2 | |
| 2068 | 34 | No | Travel_Rarely | 628 | Research & Development | 8 | |

5 rows × 34 columns




In [13]: `employee.tail(7)`

Out[13]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|----------------|-----|-----------|-------------------|-----------|------------------------|------------------|-------|
| EmployeeNumber | | | | | | | |
| 2057 | 31 | No | Non-Travel | 325 | Research & Development | 5 | |
| 2060 | 26 | No | Travel_Rarely | 1167 | Sales | 5 | |
| 2061 | 36 | No | Travel_Frequently | 884 | Research & Development | 23 | |
| 2062 | 39 | No | Travel_Rarely | 613 | Research & Development | 6 | |
| 2064 | 27 | No | Travel_Rarely | 155 | Research & Development | 4 | |
| 2065 | 49 | No | Travel_Frequently | 1023 | Sales | 2 | |
| 2068 | 34 | No | Travel_Rarely | 628 | Research & Development | 8 | |

7 rows × 34 columns



In [14]:

```
employee.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1470 entries, 1 to 2068
Data columns (total 34 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Age                                   1470 non-null   int64
 1   Attrition                           1470 non-null   object
 2   BusinessTravel                       1470 non-null   object
 3   DailyRate                           1470 non-null   int64
 4   Department                           1470 non-null   object
 5   DistanceFromHome                    1470 non-null   int64
 6   Education                           1470 non-null   int64
 7   EducationField                       1470 non-null   object
 8   EmployeeCount                       1470 non-null   int64
 9   EnvironmentSatisfaction              1470 non-null   int64
10   Gender                               1470 non-null   object
11   HourlyRate                           1470 non-null   int64
12   JobInvolvement                       1470 non-null   int64
13   JobLevel                             1470 non-null   int64
14   JobRole                              1470 non-null   object
15   JobSatisfaction                      1470 non-null   int64
16   MaritalStatus                       1470 non-null   object
17   MonthlyIncome                       1470 non-null   int64
18   MonthlyRate                         1470 non-null   int64
19   NumCompaniesWorked                  1470 non-null   int64
20   Over18                              1470 non-null   object
21   OverTime                            1470 non-null   object
22   PercentSalaryHike                   1470 non-null   int64
23   PerformanceRating                   1470 non-null   int64
24   RelationshipSatisfaction             1470 non-null   int64
25   StandardHours                       1470 non-null   int64
26   StockOptionLevel                    1470 non-null   int64
27   TotalWorkingYears                   1470 non-null   int64
28   TrainingTimesLastYear               1470 non-null   int64
29   WorkLifeBalance                     1470 non-null   int64
30   YearsAtCompany                      1470 non-null   int64
31   YearsInCurrentRole                  1470 non-null   int64
32   YearsSinceLastPromotion             1470 non-null   int64
33   YearsWithCurrManager                1470 non-null   int64
dtypes: int64(25), object(9)
memory usage: 402.0+ KB

```

Task #4: Calculate Summary Statistics for the DataFrame's Columns

1. Write code to compute summary statistics for the numeric variables.
2. Write code to compute summary statistics for just the Age column.
3. Repeat Step #1, but also include the 10th and 90th percentiles in the summary statistics that you compute.
4. Write code to compute summary statistics for the string variables.
5. Write code to compute summary statistics for just the MaritalStatus column.
6. Pick out 3 numeric variables that you think will be important in analyzing employee attrition and compute summary statistics for just those numeric variables.

```
In [15]: employee.describe()
```

Out[15]:

| | Age | DailyRate | DistanceFromHome | Education | EmployeeCount | EnvironmentSatisfacti |
|-------|-------------|-------------|------------------|-------------|---------------|-----------------------|
| count | 1470.000000 | 1470.000000 | 1470.000000 | 1470.000000 | 1470.0 | 1470.0000 |
| mean | 36.923810 | 802.485714 | 9.192517 | 2.912925 | 1.0 | 2.7217 |
| std | 9.135373 | 403.509100 | 8.106864 | 1.024165 | 0.0 | 1.0930 |
| min | 18.000000 | 102.000000 | 1.000000 | 1.000000 | 1.0 | 1.0000 |
| 25% | 30.000000 | 465.000000 | 2.000000 | 2.000000 | 1.0 | 2.0000 |
| 50% | 36.000000 | 802.000000 | 7.000000 | 3.000000 | 1.0 | 3.0000 |
| 75% | 43.000000 | 1157.000000 | 14.000000 | 4.000000 | 1.0 | 4.0000 |
| max | 60.000000 | 1499.000000 | 29.000000 | 5.000000 | 1.0 | 4.0000 |

8 rows × 25 columns

In [16]:

employee.Age.describe()

Out[16]:

count 1470.000000
mean 36.923810
std 9.135373
min 18.000000
25% 30.000000
50% 36.000000
75% 43.000000
max 60.000000
Name: Age, dtype: float64

In [17]:

employee.Age.describe(percentiles=[.1, .25,.50,.75, .9])

Out[17]:

count 1470.000000
mean 36.923810
std 9.135373
min 18.000000
10% 26.000000
25% 30.000000
50% 36.000000
75% 43.000000
90% 50.000000
max 60.000000
Name: Age, dtype: float64

In [18]:

employee.describe(include=[object])

Out[18]:

| | Attrition | BusinessTravel | Department | EducationField | Gender | JobRole | MaritalStatus | Over1 |
|--------|-----------|----------------|------------------------|----------------|--------|-----------------|---------------|-------|
| count | 1470 | 1470 | 1470 | 1470 | 1470 | 1470 | 1470 | 1470 |
| unique | 2 | 3 | 3 | 6 | 2 | 9 | 3 | 14 |
| top | No | Travel_Rarely | Research & Development | Life Sciences | Male | Sales Executive | Married | 1470 |
| freq | 1233 | 1043 | 961 | 606 | 882 | 326 | 673 | 1470 |

Task #5: Answer Simple Questions about the Dataset

The HR director asks you to answer a few descriptive questions about employees, you use this dataset to answer them: - How many employees are there by department in the dataset? - What is the overall attrition rate? - What is the average hourly rate and average yearly income? - What is the average number of years at the company? - Who are the 5 employees with the most number of years at the company? - How satisfied are employees overall?

Question 1: How many employees are there by department in the dataset?

```
In [19]: employee['Department'].value_counts()
```

```
Out[19]: Research & Development    961  
Sales                             446  
Human Resources                  63  
Name: Department, dtype: int64
```

Question 2: What is the overall attrition rate?

```
In [20]: employee['Attrition'].value_counts(normalize=True)['Yes']*100
```

```
Out[20]: 16.122448979591837
```

Question 3: What is the average hourly rate and average yearly income?

```
In [21]: print(f"Average Hourly Rate(Mean): {employee.HourlyRate.mean()}\n")  
print(f"Average Hourly Rate(Median): {employee.HourlyRate.median()}\n")  
print(f"Average Yearly Income(Mean): {employee.MonthlyIncome.sum() * 12 / len(employee.MonthlyIncome)}\n")  
print(f"Average Yearly Income(Median): {employee.MonthlyIncome.median() * 12}\n")
```

```
Average Hourly Rate(Mean): 65.89115646258503
```

```
Average Hourly Rate(Median): 66.0
```

```
Average Yearly Income(Mean): 78035.17551020408
```

```
Average Yearly Income(Median): 59028.0
```

Question 4: What is the average number of years at the company?

```
In [22]: employee.YearsAtCompany.mean()
```

```
Out[22]: 7.0081632653061225
```

Question 5: Who are the 5 employees with the most number of years at the company?

```
In [23]: employee.YearsAtCompany.sort_values(ascending=False).head()
```

```
Out[23]: EmployeeNumber
165      40
131      37
1578     36
374      36
776      34
Name: YearsAtCompany, dtype: int64
```

```
In [24]: employee.YearsAtCompany.sort_values(ascending=False).value_counts()
```

```
Out[24]: 5      196
1      171
3      128
2      127
10     120
4      110
7       90
9       82
8       80
6       76
0       44
11      32
20      27
13      24
15      20
14      18
22      15
12      14
21      14
18      13
16      12
19      11
17       9
24       6
33       5
25       4
26       4
31       3
32       3
36       2
27       2
29       2
23       2
30       1
34       1
37       1
40       1
Name: YearsAtCompany, dtype: int64
```

Question 6: How satisfied are employees overall?

```
In [25]: employee.JobSatisfaction.value_counts()
```

```
Out[25]: 4    459
          3    442
          1    289
          2    280
          Name: JobSatisfaction, dtype: int64
```


```
In [26]: JobSatisfaction_cat = {
          1: 'Low',
          2: 'Medium',
          3: 'High',
          4: 'Very High'
        }
employee.JobSatisfaction = employee.JobSatisfaction.map(JobSatisfaction_cat)
```

```
In [27]: 100 * employee.JobSatisfaction.value_counts(normalize=True)
```

```
Out[27]: Very High    31.224490
          High        30.068027
          Low         19.659864
          Medium      19.047619
          Name: JobSatisfaction, dtype: float64
```

```
In [28]: employee.JobSatisfaction.value_counts().plot(kind = 'bar')
```

```
Out[28]: <AxesSubplot:>
```



Task #6: Answer Additional Questions about the Dataset

After taking a look at your answers the HR director, asks you more questions: - Give me the list of the employees with Low level of JobSatisfaction - Give me the list of the employees with Low level of both JobSatisfaction and PerformanceRating - Compare the employees with Low and Very High JobSatisfaction across the following variables: Age, Department, - DistanceFromHome, HourlyRate, MonthlyIncome and YearsAtCompany.

Question 7: Give me the list of the employees with Low level of JobSatisfaction

```
In [29]: employee.loc[employee.JobSatisfaction == 'Low'].index
```

```
Out[29]: Int64Index([ 10,  20,  27,  31,  33,  38,  51,  52,  54,  68,  
                  ...  
                  1975, 1980, 1998, 2021, 2023, 2038, 2054, 2055, 2057, 2062],  
                  dtype='int64', name='EmployeeNumber', length=289)
```

```
In [30]: employee[employee.JobSatisfaction == 'Low'].JobSatisfaction
```

```
Out[30]: EmployeeNumber  
10      Low  
20      Low  
27      Low  
31      Low  
33      Low  
...  
2038    Low  
2054    Low  
2055    Low  
2057    Low  
2062    Low  
Name: JobSatisfaction, Length: 289, dtype: object
```

Question 8: Give me the list of the employees with Low level of both JobSatisfaction and JobInvolment

```
In [31]: employee.JobInvolvement.value_counts()
```

```
Out[31]: 3      868  
         2      375  
         4      144  
         1       83  
         Name: JobInvolvement, dtype: int64
```

```
In [32]: employee.JobInvolvement = employee.JobInvolvement.map(JobSatisfaction_cat)
```

```
In [33]: employee.loc[(employee.JobSatisfaction == 'Low') & (employee.JobInvolvement == 'Low')].index
```

```
Out[33]: Int64Index([33, 235, 454, 615, 1019, 1037, 1237, 1460, 1478, 1544, 1611, 1622,  
                  1905, 1956],  
                  dtype='int64', name='EmployeeNumber')
```

Question 9: Compare the employees with Low and Very High JobSatisfaction across the following variables: Age, Department, DistanceFromHome, HourlyRate, MonthlyIncome and YearsAtCompany.

```
In [38]: subset_of_interest = employee.loc[(employee.JobSatisfaction == 'Low') | (employee.JobSa
```

```
In [39]: subset_of_interest.shape
```

```
Out[39]: (748, 34)
```

```
In [41]: subset_of_interest.JobSatisfaction.value_counts()
```

```
Out[41]: Very High    459
         Low         289
         Name: JobSatisfaction, dtype: int64
```

```
In [42]: grouped = subset_of_interest.groupby('JobSatisfaction')
```

```
In [50]: grouped.groups
```

```
Out[50]: {'Low': [10, 20, 27, 31, 33, 38, 51, 52, 54, 68, 70, 74, 75, 81, 86, 88, 100, 101, 113, 124, 133, 134, 145, 153, 170, 190, 197, 199, 200, 235, 239, 240, 241, 244, 250, 267, 274, 282, 288, 297, 299, 303, 328, 334, 339, 340, 347, 351, 362, 369, 374, 382, 390, 396, 412, 424, 425, 429, 451, 454, 474, 486, 510, 515, 517, 522, 524, 530, 532, 534, 536, 538, 549, 567, 573, 590, 605, 615, 625, 630, 648, 650, 662, 664, 667, 682, 684, 702, 705, 725, 728, 729, 732, 733, 742, 758, 764, 771, 775, 776, ...], 'Very High': [1, 8, 18, 22, 23, 24, 30, 36, 39, 40, 42, 45, 49, 53, 57, 62, 63, 72, 73, 76, 78, 79, 97, 98, 104, 106, 107, 112, 116, 117, 118, 120, 137, 139, 140, 143, 144, 148, 152, 154, 155, 158, 165, 169, 174, 179, 184, 192, 195, 198, 207, 215, 217, 221, 223, 228, 230, 242, 243, 245, 246, 262, 264, 273, 275, 281, 283, 286, 287, 291, 298, 302, 306, 309, 311, 312, 315, 316, 319, 323, 325, 327, 333, 335, 336, 338, 346, 349, 353, 361, 367, 372, 373, 377, 378, 380, 388, 389, 391, 393, ...]}
```

```
In [52]: grouped.get_group('Low').head()
```

```
Out[52]:
```

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|----------------|-----|-----------|-------------------|-----------|------------------------|------------------|-------|
| EmployeeNumber | | | | | | | |
| 10 | 59 | No | Travel_Rarely | 1324 | Research & Development | | 3 |
| 20 | 29 | No | Travel_Rarely | 1389 | Research & Development | | 21 |
| 27 | 36 | Yes | Travel_Rarely | 1218 | Sales | | 9 |
| 31 | 34 | Yes | Travel_Rarely | 699 | Research & Development | | 6 |
| 33 | 32 | Yes | Travel_Frequently | 1125 | Research & Development | | 16 |

5 rows × 34 columns


```
In [54]: grouped.get_group('Very High').head()
```

Out[54]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Educa |
|--|-----|-----------|----------------|-----------|------------|------------------|-------|
|--|-----|-----------|----------------|-----------|------------|------------------|-------|

| EmployeeNumber | | | | | | | |
|----------------|----|-----|-------------------|------|------------------------|--|----|
| 1 | 41 | Yes | Travel_Rarely | 1102 | Sales | | 1 |
| 8 | 32 | No | Travel_Frequently | 1005 | Research & Development | | 2 |
| 18 | 34 | No | Travel_Rarely | 1346 | Research & Development | | 19 |
| 22 | 22 | No | Non-Travel | 1123 | Research & Development | | 16 |
| 23 | 53 | No | Travel_Rarely | 1219 | Sales | | 2 |

5 rows × 34 columns



Age

In [56]: `grouped['Age'].mean()`

Out[56]:

| JobSatisfaction | |
|-----------------|-----------|
| Low | 36.916955 |
| Very High | 36.795207 |

Name: Age, dtype: float64

In [58]: `grouped['Age'].describe().unstack()`

Out[58]:

| JobSatisfaction | | |
|-----------------|-----------|------------|
| count | Low | 289.000000 |
| | Very High | 459.000000 |
| mean | Low | 36.916955 |
| | Very High | 36.795207 |
| std | Low | 9.245496 |
| | Very High | 9.125609 |
| min | Low | 19.000000 |
| | Very High | 18.000000 |
| 25% | Low | 30.000000 |
| | Very High | 30.000000 |
| 50% | Low | 36.000000 |
| | Very High | 35.000000 |
| 75% | Low | 42.000000 |
| | Very High | 43.000000 |
| max | Low | 60.000000 |
| | Very High | 60.000000 |

dtype: float64

Department

In [59]: `grouped['Department'].describe()`

Out[59]:

| | count | unique | top | freq |
|--|-------|--------|-----|------|
|--|-------|--------|-----|------|

JobSatisfaction

| | | | | |
|-----------|-----|---|------------------------|-----|
| Low | 289 | 3 | Research & Development | 192 |
| Very High | 459 | 3 | Research & Development | 295 |

DistanceFromHome

In []:

HourlyRate

In []:

MonthlyIncome

In []:

YearsAtCompany

In []:

Task #7: Create a DataFrame to Compare the Means Across All Numerical Variables

Comparing the means across all numerical variables

Although we we asked for just some specific columns, to give the HR director a better picture of how these groups compare across different variables, let's create a DataFrame that contains the mean for every numeric variable in our dataset.

In []:

Task #8: Create Some Additional Questions That You Could Ask About the Dataset

In []:

Final Note:

Issues to keep in mind about this dataset:

- There are many variables that are detected as numerical but are actually categorical (like Education).
- Since this is a simulated dataset, it is hard to find interesting patterns.

